


云原生社区Meetup第一期



主办方  云原生社区
Cloud Native Community

承办方  UCloud 优刻得  云原生社区上海站
Cloud Native Community Shanghai

赞助商  tetrate  万科海外  Broadview

2020-11-28 上海站



高洪涛



Tetrate 创始工程师

原华为软件开发云技术专家，对云APM有深入的理解，并有丰富的APM产品设计，研发与实施经验。对分布式数据库，容器调度，微服务，ServiceMesh等技术有深入的了解。

云原生社区 Meetup
第一期 · 上海站



Apache SkyWalking Adapter for HPA

演讲人：高洪涛

gaohontao Tetrade

▶ 目录

章节一 Apache SkyWalking

章节二 Kubernetes HPA

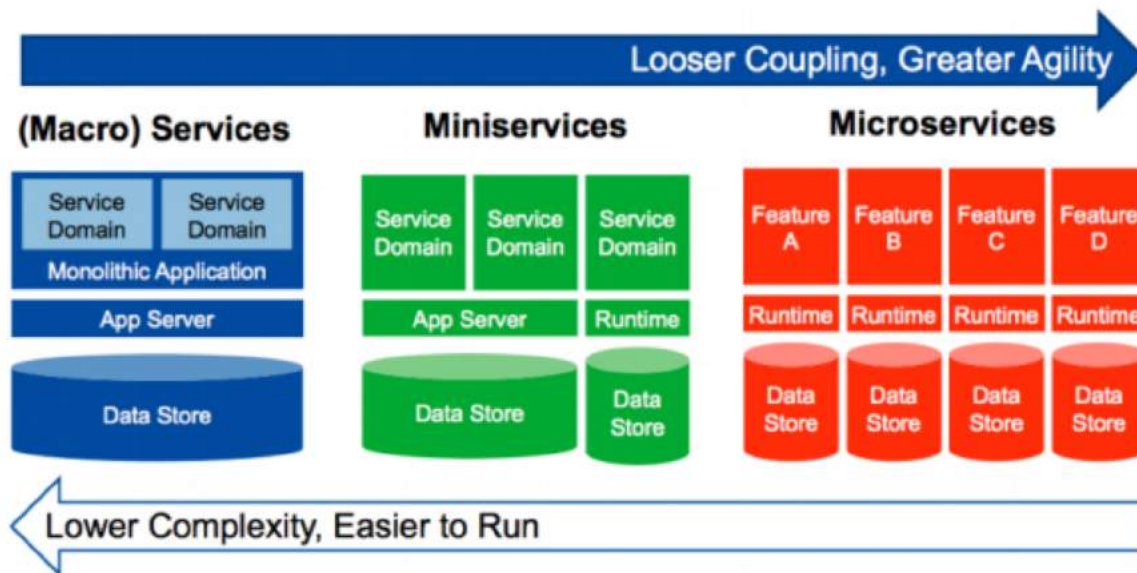
章节三 Apache SkyWalking Adapter

Apache SkyWalking



Service in real world


Think Multigrained, Not Just "Micro"



20 © 2018 Gartner, Inc. and/or its affiliates. All rights reserved.

Gartner.

Independent
Loosely coupled
Multiple layers



O11Y

Logging


By Tracing
Satellite

Metrics

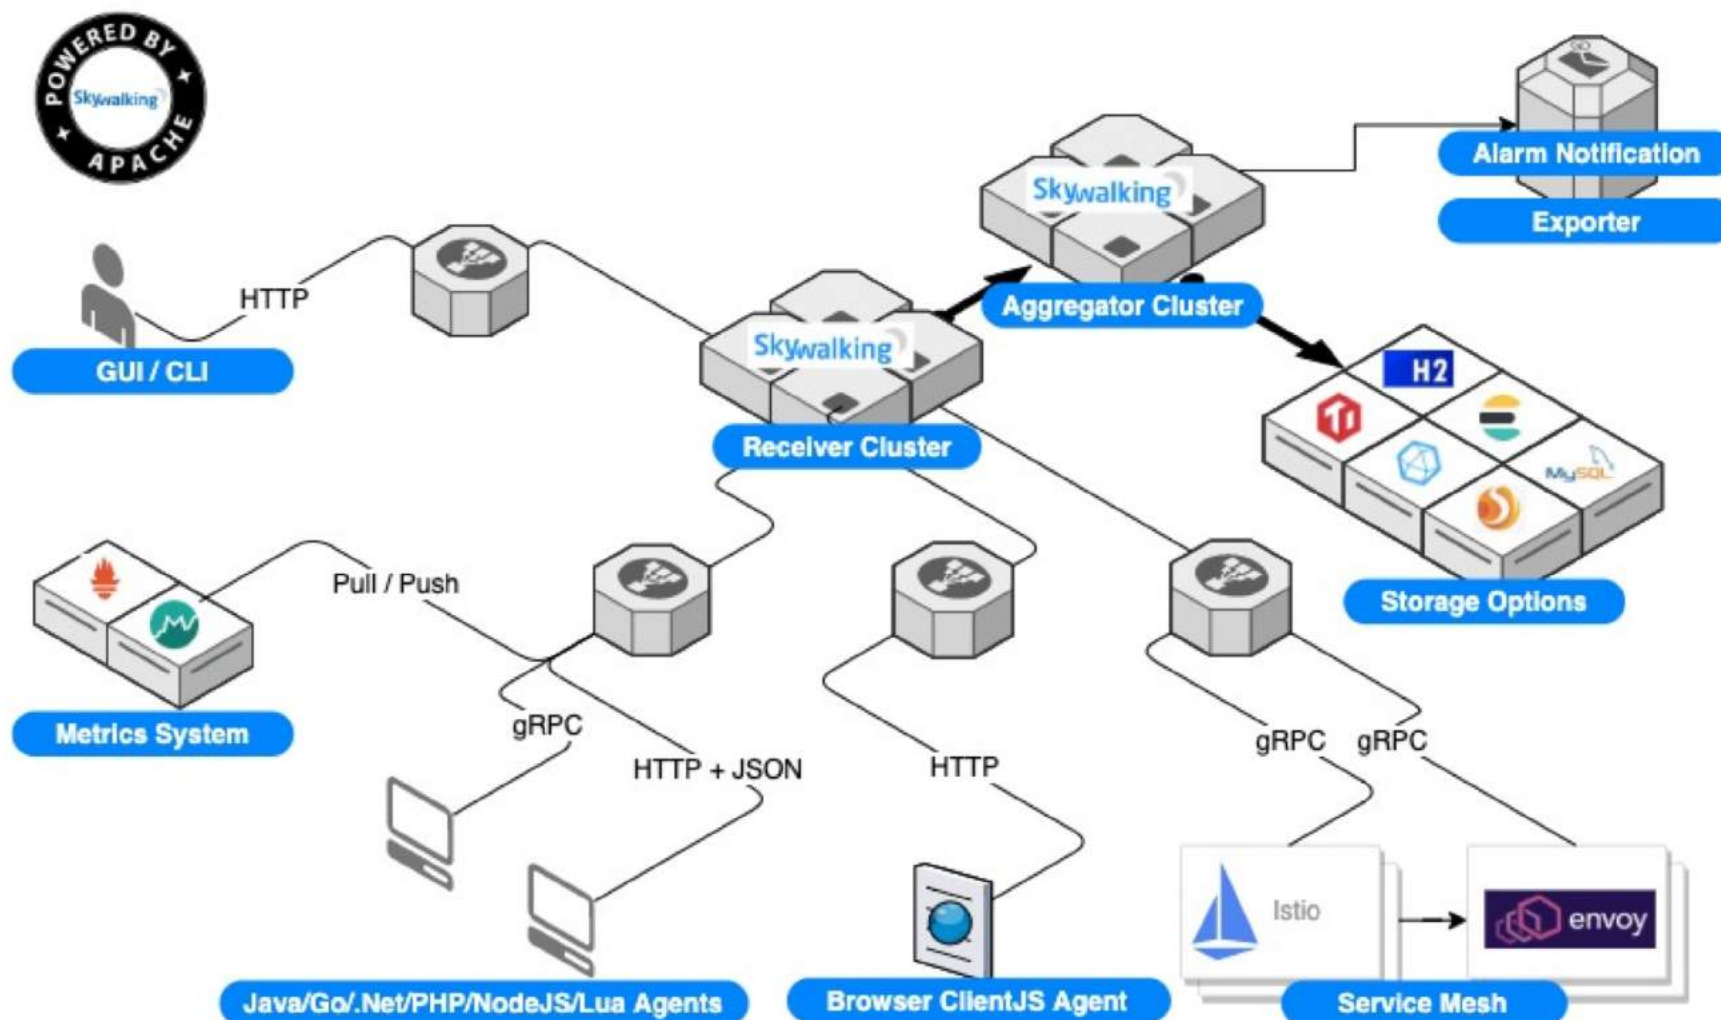
Meter & MAL
system


Tracing

SW Native Agent
Nginx & Envoy



Arch






Cloud on Kubernetes


Not just running, but living with Kubernetes



Operator




K8S O11Y



HPA, VPA



Istio



User
friendly



More...

Query Interface

SWCTL

User friendly


Imported by Go

CLI

UI

GraphQL

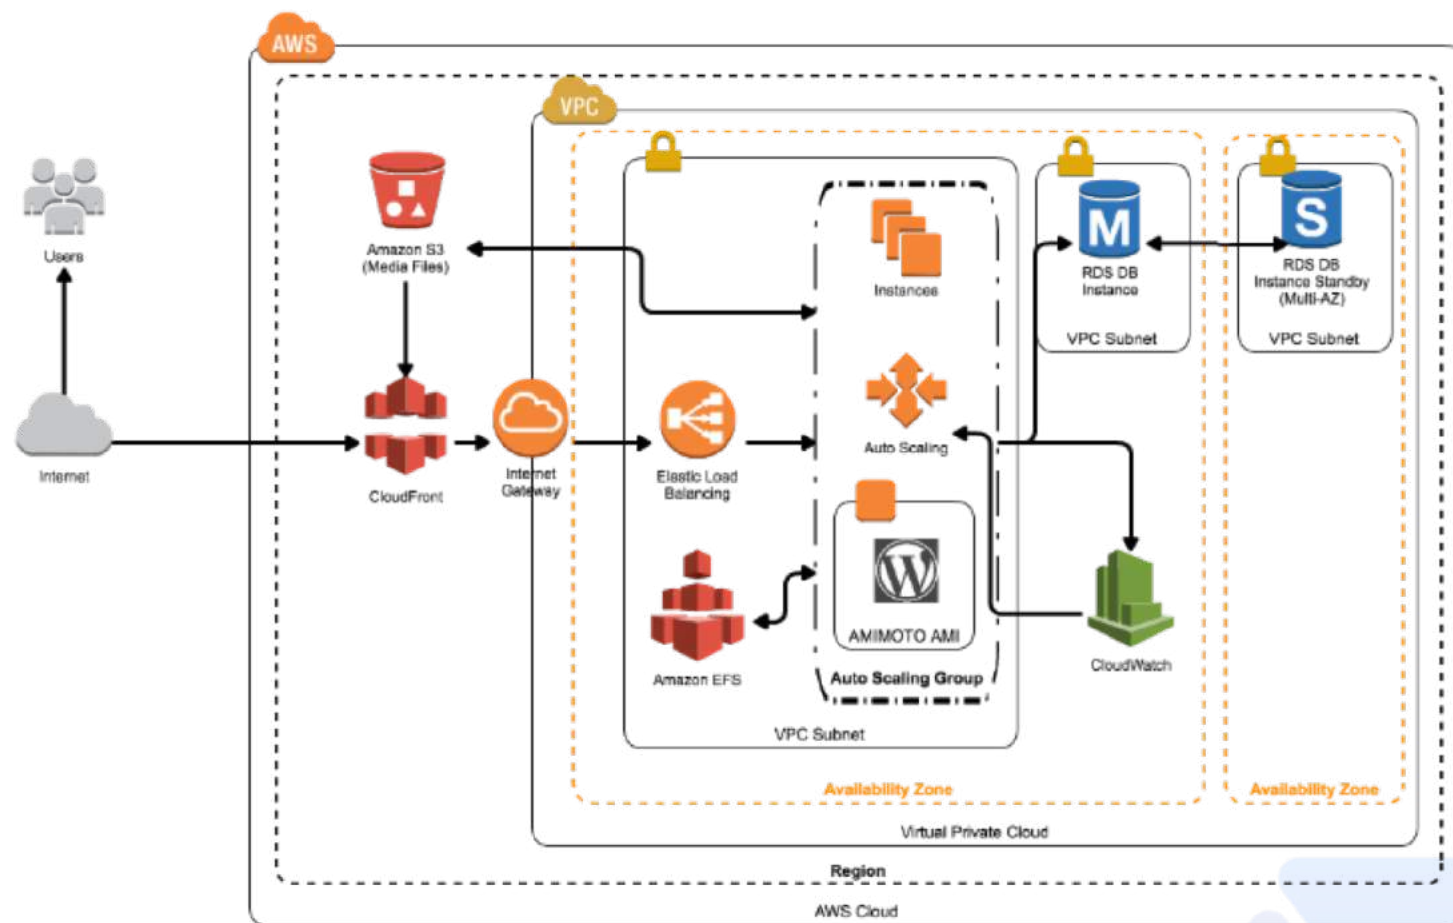
Aggregation
Strongly typed



Kubernetes HPA



VM Auto-scaling



What's auto scaling

Autoscaling on an abstract level:


- Calculate resources to cover demand
- Demand measured by metrics
 - Metrics must be collected, stored and made queriable
 - fulfill SLO



Horizontal vs Vertical


HAP:

Increase replicas when necessary



VPA:


Increase resource request/limit when necessary



Problem of HPA v1

Loosely specified APIs


Unmaintained hipster



Problem of HPA v1


Loosely specified APIs

Unmaintained hipster



Resource and Custom metrics API

- An API specification
- Implemented and maintained by vendors
- Returns single value



Types of API

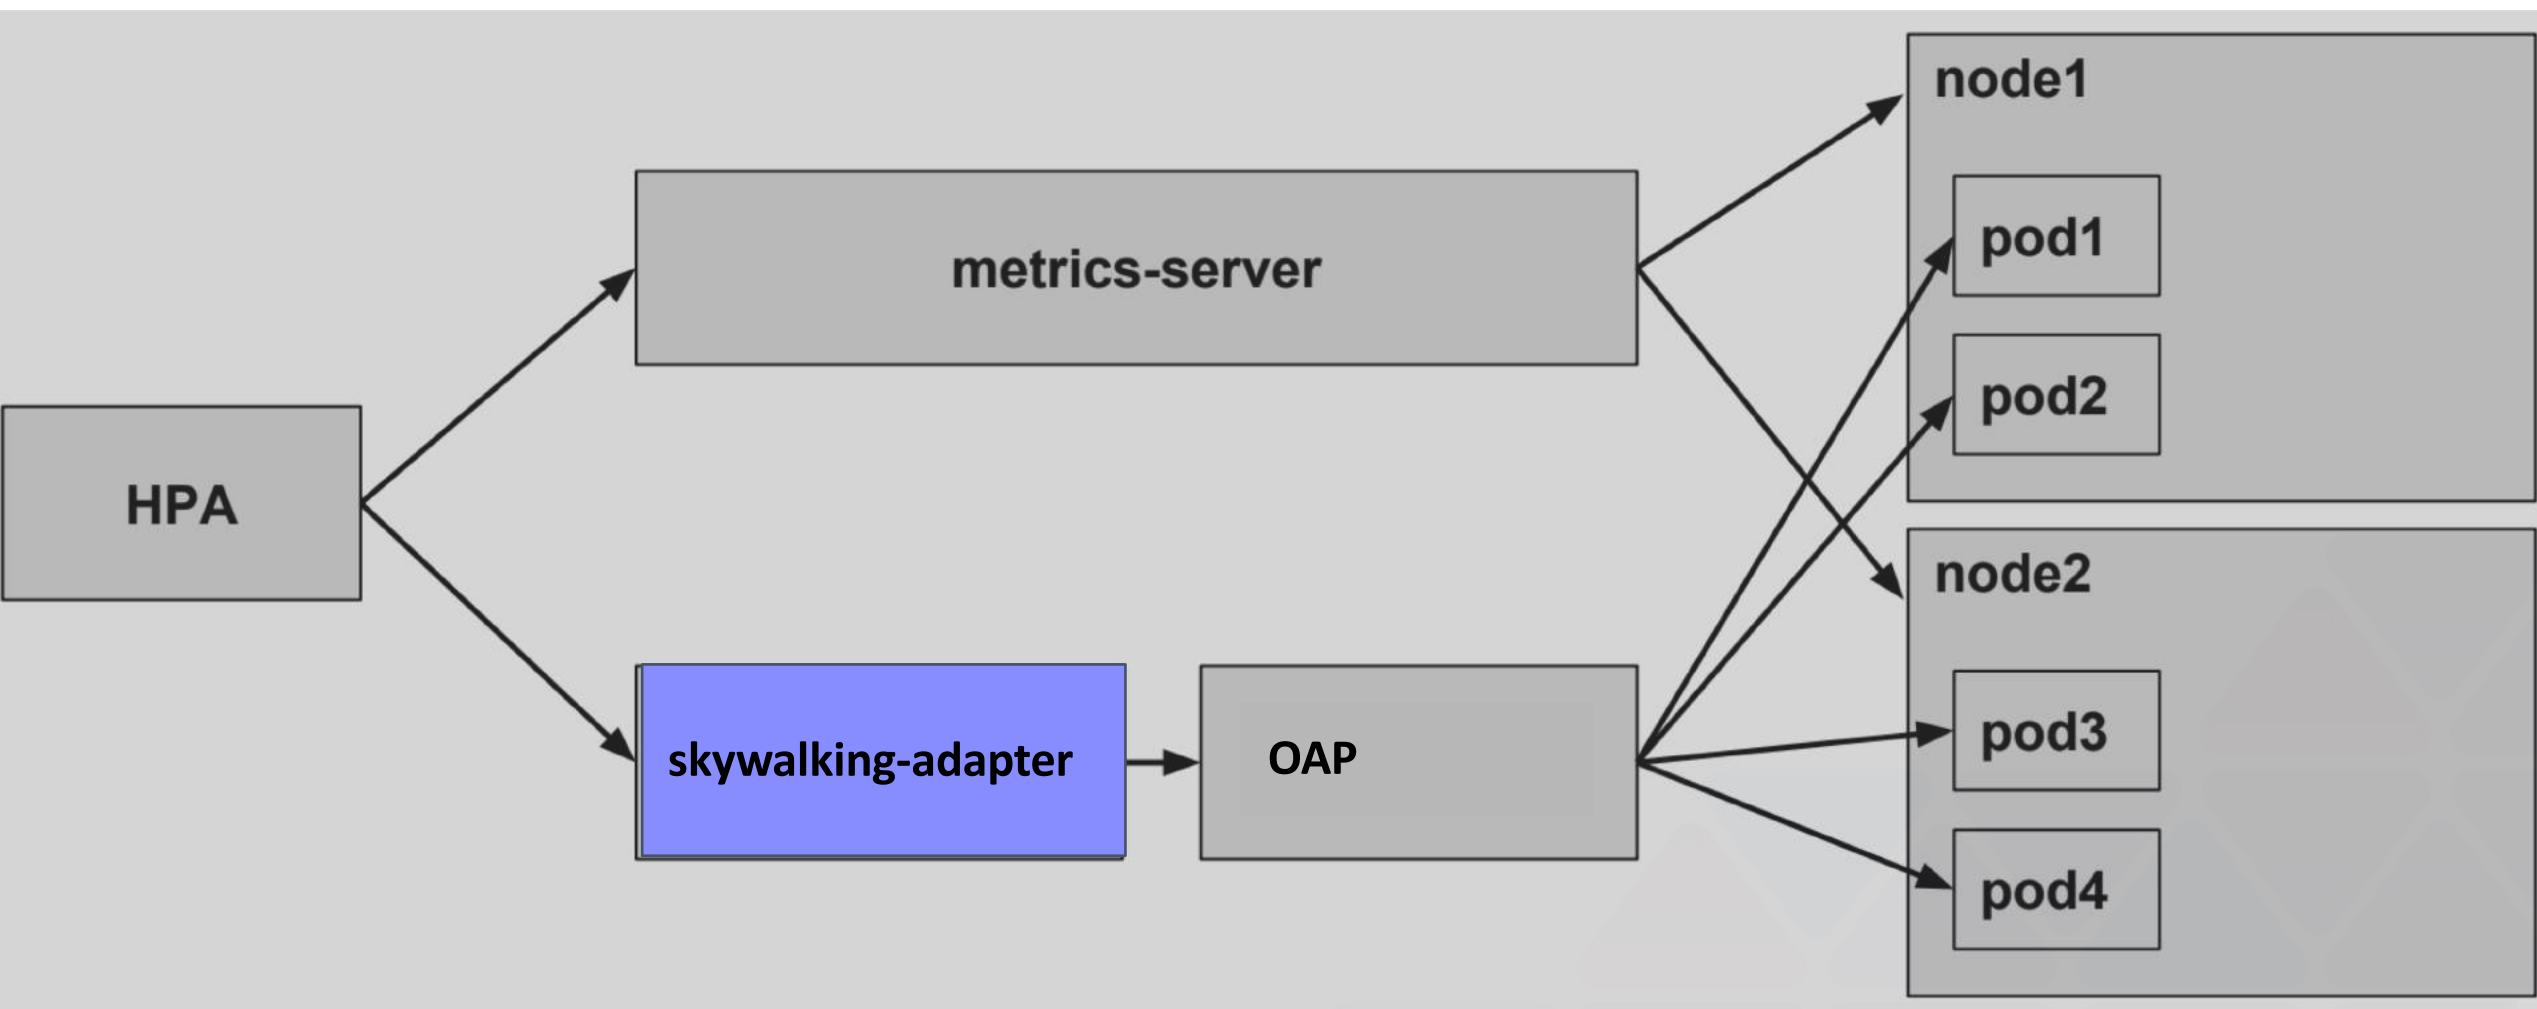
- Resource Metrics API
- Custom Metrics API
- External Metrics API



Apache SkyWalking Adapter



Adatper



Metric Scope

Kubernetes	SkyWalking
Service/Deployment	Service
Pod	Instance

- Deploy SkyWalking backend cluster
- Deploy adapter

<http://github.com/apache/skywalking-swck>




Query type

- Linear: single dimension, throughput, latency
- Multi-linear: multiple dimension, histogram, percentile



Don't query the latest

- The value varies
- Affect the behavior of HPA



Demo

```
- type: External
  external:
    metric:
      name: service_percentile
      metricSelector:
        matchLabels:
          service: front_gateway
          # The index of [P50, P75, P90, P95, P99]. 2 is the index of P90(90%)
          label: "2"
    target:
      type: value
      value: 80
```



云原生社区Meetup

第一期·上海站



Wechat: putao209764

Twitter: @hanahmily

Mail: hanahmily@gmail.com